Programming C in a UNIX Environment – Review

GMU Fall 2019 CS 367

Connor Baker

Compiled on August 28, 2019 at 10:52am

Reading Track

Reading homework:

1. Read through Chapter 2.1 (bits and information storage) of the textbook
    - The next lecture begins with bitwise, logical, and shift operations in C
2. Continue to review C. If you know these well in C, it will be a lot easier over the next two months.
    - Passing by reference and working with pointers in functions
    - Dereferencing values: `*ptr = 42;`
    - Pointer arithmetic with arrays: `*(ary + 3)= 42;`
    - Working with `struct`s and dynamic memory
    - Working with linked lists using dynamically allocated `struct`s

Lecutre Overview from CS262

Describe the Memory Representation of a Pointer	Relate Pointers to Addresses	K&R 5.1
Indirectly Access Primitive Data Types in C	Use Pointers in C to Read/Write Values in Memory	
Declare and Access Data from Arrays in C	Declare Single and Multi-Dimensional Arrays in C	K&R 1.6
	Access Elements of Single and Multi-Dimensional Arrays in C	
Evaluate Array References in C for Correctness	Determine Out-of-Bounds Errors in Provided C Statements	
Pass Arrays to Functions in C	Pass an Array by Name to a Function as an Argument in C	K&R 5.3
Indirectly Reference Array Elements in C	Use Pointers in C to Read/Write Array Elements	
	Use Pointer Arithmetic and an Array Name to Access Elements	
List the Three Standard I/O Streams	Identify the Three Standard I/O Streams	K&R 7.5
Perform Common Standard I/O Operations in C	Open, Read, Write, and Close Files in C	
	Get User Input using stdin	
	Format User Output using stdout/stderr	
Describe the Debugging Process	List Debugging Goals	K&R 7.6
	List Programming Errors	
	Describe Ad-Hoc Debugging	
	Describe Step-by-step (process) Debugging	
Define and Allocate Memory for Structs in C	Define a Struct in C	K&R 6.1
	Declare a Struct in C using Static Memory	
	Declare a Struct in C using Dynamic Memory	K&R 6.4
Access Struct Elements in C	Access Elements of a Static Struct with the . notation	K&R 6.1
	Access Elements of a Dynamic Struct with the -> notation	K&R 6.2
	Access Elements of a Struct using Pointers and Pointer Arithmetic	K&R 6.1
Create Custom Types in C	Use typedef to Create a Custom Type in C	K&R 6.7
Dynamically Allocate and Release Data in C	Use malloc/free to create pointers in C	K&R 6.5
	Describe the Heap and the Arbitrariness of Allocation Locations	
Create and Access Data from a Linked List in C	Define a Struct in C Containing a Pointer to the Struct Type	
	Declare and Add Structs to the Linked List in a Given Manner	
	Find Given Data in a Linked List in O(n) time.	

Figure 1: Lecture overview.

Programming Tools in UNIX

- Integrated Development Environments (IDE)

  - Visual Studio, Netbeans, Eclipse, jGRASP
  - These don't work (or work well) over a remote terminal
    * In Systems Programming, it's common to remotely connect to an embedded device and work directly on it
    * This semester, we'll be working on the Zeus server remotely

- You don't need an IDE. All you need are
  - A text editor (`vi`/`vim` or `emacs`)
  - A compiler (`gcc`)
  - A debugger (`gdb`)

Compiling on Zeus

The general template we'll be using is

```
gcc -g -01 -o executable source_file.c
```

The flags used in this example are

- `-g` Compile with symbols in the code. Useful for debugging.
- `-O1` Compile with optimization level one. This reduces the code/memory footprint.
- `-o` Specify the output filename. If you don't specify the output, you'll get `a.out`.

Makefiles

Makefiles are super useful for building multi-source file programs.

There are three parts to each *rule* in a Makefile. Let's diagram walk through the following rule.

```
prog: prog.c lib.c
    gcc -g -o prog prog.c lib.c
```

Component	Example
Rule name	`prog`
Dependencies	`prog.c lib.c`
Action	`gcc -g -o prog prog.c lib.c`

The Makefile ensures that the dependencies have been satisfied before running the action. If you were to put another rule as a dependency, it would verify that *that* dependency was satisfied before continuing.

Note that the tab before the action is important!

Pointers

At the machine level, you get two types of data: *values* and *addresses*.

C pointers let us work with addresses.

Before we continue, let's look at the size of different data types on Zeus.

Data Type	Size
`char`	1 B
`short`	2 B
`int`	4 B
`long`	8 B
`char *`	8 B
`short *`	8 B
`int *`	8 B
`long *`	8 B

Example 1

Consider the following example:

```
1  int main() {
2    int a = 42;
3    int *p;
4  }
```

1. How do we set `p` to point to `a`?
    - Let `p` reference the memory address of `a`: `p = &a;`
2. How do we set `a` to 12 using `p`?
    - Dereference `p` to change it's value: `*p = 12;`

Question 1

File: `ptr_2.c`

What will the output of the following program be?

```c
#include <stdio.h>

int main() {
  int a = 42, b = 16;
  int *p, *q;
  p = &a;
  q = &b;
  *q = *p;
  p = q;
  printf("%d %d %d %d\n", a, b, *p, *q);
}
```

The output is:

```
$ ./ptr_2
42 42 42 42
```

The pointee of `q` is set equal to the pointee of `p` with the statement `*q = *p;`. The statement `p = q;` changes `p` so that it points at the same thing that `q` does – a value which was set equal the thing that `p` *was* pointing at earlier.

Question 2

Write a function called swap to exchange two integers by reference.

```c
void swap(int *a, int *b) {
  int temp = *a;
  *a = *b;
  *b = temp;
}
```

Arrays

An array is a list of values arranged consecutively in memory.

Arrays in C use the `[]` bracket notation which is indexed by offsets. In fact, `a[n]` is really just syntactic sugar for `*(a + n)`. Remember that pointers are a different type: when you add an integer to a pointer, it's scaled by something akin to the `sizeof` function.

Some examples:

- `int a[5];` declares an array of five `int`s (on Zeus this is 20 bytes)
- `a[0]` references the value at address `a`
- `a[1]` references the value at the address `(a+1)`.

  - As noted above, in hardware this would be the address of `a` plus four bytes (since an `int` on Zeus is four bytes)

- `a[4]` references the last element of `a`
- `a[5]` is out of bounds but completely legal C – no bounds checking, remember?

Arrays and Pointers

Files: `ary_a.c` and `ary_b.c`.

The name of the array is the address at of the start of the array. An array name is *not* a variable – it is a constant value.

This may be surprising. "But an array name is an `lvalue`!" you might claim – and you'd be correct. However, an array name is not a *modifiable* `lvaule`. For an excellent breakdown, see this Stack Overflow post:

Is there a reason why an array name is not an lvalue?

Pointer Arithmetic in Arrays

File: `ptr_ary.c`.

You can use pointer arithmetic when working with arrays.

```
1  int nums[5] = {1, 6, 10, 42, -14};
2  int *p_nums = NULL;
3  p_nums = nums;
4  // Address of the array
5  printf("Address of nums: %p\n", nums);
6  printf("Address p_nums points to: %p\n", p_nums);
7  // Access the fourth number with [] notation
8  printf("p_nums[3] == %d\n", p_nums[3]);
9  // Access the fourth number with pointer arithmetic
10 printf("p_nums[3] == %d\n", *(p_nums + 3));
```

Question 3

Given the following snippet:

```
1  int nums[5] = {1, 6, 10, 42, -14};
2  int *p_nums = &nums[3];
3  p_nums = nums;
4  // Address of nums = 0x400
5  // sizeof(int) = 4 bytes, sizeof(int *) = 8 bytes
```

what is the value of each of the following expressions?

1. `&nums[2]`
    - 0x408
2. `p_nums[-1]`
    - 10
3. `*(p_nums + 1)`
    - -14

I/O in C

What is `stdio.h`?

- The header file for the C Standard Library (`libc.a`)

Okay, so what's `libc.a`?

- A static library which is automatically included
- It's a set of pre-compiled C objects which contain all of the core C functions
- `stdio.h` only has the prototypes and macros for those functions

Some examples of standard I/O operations:

- Opening and closing files: `fopen` and `fclose`
- Reading and writing files: `fread` and `fwrite`
- Reading and writing text: `fgets` and `fputs`
- Formatted reading and writing: `fscanf` and `fprintf`

Standard I/O is implemented on top of the Operating System (OS) I/O:

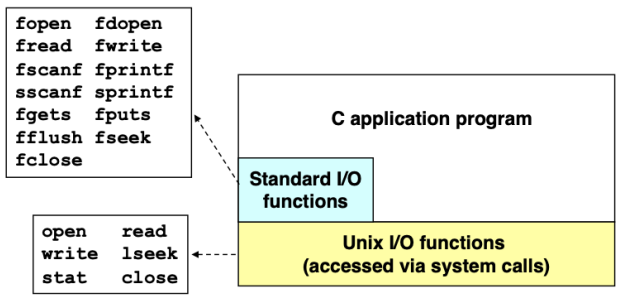

Figure 2: Unix I/O interacting with C I/O.

Standard I/O Streams in C and Unix

Standard I/O models all open files as streams. This is an abstraction for a file descriptor and a buffer.

There are three streams all C programs start with access to:

Name	Type	Buffered?	File Descriptor
`stdin`	Standard Input	yes	0
`stdout`	Standard Output	yes	1
`stderr`	Standard Error	no	2

Common I/O Functions

These are all declared within the `stdio.h` header.

Function	Description
`putchar`	Displays an ASCII character to the screen
`getchar`	Reads an ASCII character from the keyboard
`printf`	Displays a formatted string
`scanf`	Reads a formatted string
`fopen`	Open or create a file for I/O
`fprintf`	Writes a formatted string to a file
`fscanf`	Reads a formatted string from a file

Formatted I/O

Both `printf` and `scanf` allow conversion between ASCII representations and internal data types.

The format string of these functions contains text to be read/written as well as the format characters which describe their formatting.

Code	Description
`%d`	Signed decimal integer
`%f`	Signed decimal floating-point number

Code	Description
`%x`	Hexadecimal number
`%e`	Scientific notation
`%c`	ASCII character
`%s`	ASCII string
`%p`	Memory address (hex)

Both `printf` and `scanf` allow special characters as well.

Code	Description
`\n`	Newline
`\t`	Tab
`\b`	Backspace
`\\`	Backslash
`\'`	Single quote
`\"`	Double quote
`\0nnn`	ASCII code nnn in octal
`\xnnn`	ASCII code nnn in hex

printf

File: `print.c`.

Prints its first argument (the format string) to `stdout` with all of the formatting characters replaced by the ASCII representation of the corresponding arguments.

As an example, the snippet

```
int a = 100;
char c = 'z';
char hw[13] = "Hello World!";
double pi = 3.14159;
printf("a: %d, \'%c\', 0x%x\n", a, a, a);
printf("c: %d, '%c', 0x%x\n", c, c, c); printf("hw: \"%s\"\n", hw);
printf("pi: %lf, %.2lf, %e\n", pi, pi, pi);
```

prints

```
$ ./print
a: 100, 'd', 0x64
c: 122, 'z', 0x7a
hw: "Hello World!"
pi: 3.141590, 3.14, 3.141590e+00
```

when executed.