
Programming C in a UNIX Environment – Review

GMU Fall 2019 CS 367

Connor Baker



Compiled on August 28, 2019 at 11:38am

Reading Track

Reading homework:

1. Read through Chapter 2.1 (bits and information storage) of the textbook
 - The next lecture begins with bitwise, logical, and shift operations in C
2. Continue to review C. If you know these well in C, it will be a lot easier over the next two months.
 - Passing by reference and working with pointers in functions
 - Dereferencing values: `*ptr = 42;`
 - Pointer arithmetic with arrays: `*(ary + 3) = 42;`
 - Working with `structs` and dynamic memory
 - Working with linked lists using dynamically allocated `structs`

Lecture Overview from CS262

Describe the Memory Representation of a Pointer	Relate Pointers to Addresses	K&R 5.1
Indirectly Access Primitive Data Types in C	Use Pointers in C to Read/Write Values in Memory	
Declare and Access Data from Arrays in C	Declare Single and Multi-Dimensional Arrays in C	K&R 1.6
	Access Elements of Single and Multi-Dimensional Arrays in C	
Evaluate Array References in C for Correctness	Determine Out-of-Bounds Errors in Provided C Statements	
Pass Arrays to Functions in C	Pass an Array by Name to a Function as an Argument in C	K&R 5.3
Indirectly Reference Array Elements in C	Use Pointers in C to Read/Write Array Elements	
	Use Pointer Arithmetic and an Array Name to Access Elements	
List the Three Standard I/O Streams	Identify the Three Standard I/O Streams	K&R 7.5
Perform Common Standard I/O Operations in C	Open, Read, Write, and Close Files in C	
	Get User Input using <code>stdin</code>	
	Format User Output using <code>stdout/stderr</code>	
Describe the Debugging Process	List Debugging Goals	K&R 7.6
	List Programming Errors	
	Describe Ad-Hoc Debugging	
	Describe Step-by-step (process) Debugging	
Define and Allocate Memory for Structs in C	Define a Struct in C	K&R 6.1
	Declare a Struct in C using Static Memory	
	Declare a Struct in C using Dynamic Memory	K&R 6.4
Access Struct Elements in C	Access Elements of a Static Struct with the <code>.</code> notation	K&R 6.1
	Access Elements of a Dynamic Struct with the <code>-></code> notation	K&R 6.2
	Access Elements of a Struct using Pointers and Pointer Arithmetic	K&R 6.1
Create Custom Types in C	Use <code>typedef</code> to Create a Custom Type in C	K&R 6.7
Dynamically Allocate and Release Data in C	Use <code>malloc/free</code> to create pointers in C	K&R 6.5
	Describe the Heap and the Arbitrariness of Allocation Locations	
Create and Access Data from a Linked List in C	Define a Struct in C Containing a Pointer to the Struct Type	
	Declare and Add Structs to the Linked List in a Given Manner	
	Find Given Data in a Linked List in $O(n)$ time.	

Figure 1: Lecture overview.

Programming Tools in UNIX

- Integrated Development Environments (IDE)

- Visual Studio, Netbeans, Eclipse, jGRASP
- These don't work (or work well) over a remote terminal
 - * In Systems Programming, it's common to remotely connect to an embedded device and work directly on it
 - * This semester, we'll be working on the Zeus server remotely
- You don't need an IDE. All you need are
 - A text editor (`vi/vim` or `emacs`)
 - A compiler (`gcc`)
 - A debugger (`gdb`)

Compiling on Zeus

The general template we'll be using is

```
1 gcc -g -O1 -o executable source_file.c
```

The flags used in this example are

- `-g` Compile with symbols in the code. Useful for debugging.
- `-O1` Compile with optimization level one. This reduces the code/memory footprint.
- `-o` Specify the output filename. If you don't specify the output, you'll get `a.out`.

Makefiles

Makefiles are super useful for building multi-source file programs.

There are three parts to each *rule* in a Makefile. Let's diagram walk through the following rule.

```
1 prog: prog.c lib.c
2     gcc -g -o prog prog.c lib.c
```

Component	Example
Rule name	<code>prog</code>
Dependencies	<code>prog.c lib.c</code>
Action	<code>gcc -g -o prog prog.c lib.c</code>

The Makefile ensures that the dependencies have been satisfied before running the action. If you were to put another rule as a dependency, it would verify that *that* dependency was satisfied before continuing.

Note that the tab before the action is important!

Pointers

At the machine level, you get two types of data: *values* and *addresses*.

C pointers let us work with addresses.

Before we continue, let's look at the size of different data types on Zeus.

Data Type	Size
<code>char</code>	1 B
<code>short</code>	2 B
<code>int</code>	4 B
<code>long</code>	8 B
<code>char *</code>	8 B
<code>short *</code>	8 B
<code>int *</code>	8 B
<code>long *</code>	8 B

Example 1

Consider the following example:

```
1 int main() {  
2   int a = 42;  
3   int *p;  
4 }
```

1. How do we set `p` to point to `a`?
 - Let `p` reference the memory address of `a`: `p = &a;`
2. How do we set `a` to 12 using `p`?
 - Dereference `p` to change it's value: `*p = 12;`

Question 1

File: `ptr_2.c`

What will the output of the following program be?

```
1  #include <stdio.h>
2
3  int main() {
4      int a = 42, b = 16;
5      int *p, *q;
6      p = &a;
7      q = &b;
8      *q = *p;
9      p = q;
10     printf("%d %d %d %d\n", a, b, *p, *q);
11 }
```

The output is:

```
1  $ ./ptr_2
2  42 42 42 42
```

The pointee of `q` is set equal to the pointee of `p` with the statement `*q = *p;`. The statement `p = q;` changes `p` so that it points at the same thing that `q` does – a value which was set equal the thing that `p` was pointing at earlier.

Question 2

Write a function called `swap` to exchange two integers by reference.

```
1  void swap(int *a, int *b) {
2      int temp = *a;
3      *a = *b;
4      *b = temp;
5  }
```

Arrays

An array is a list of values arranged consecutively in memory.

Arrays in C use the `[]` bracket notation which is indexed by offsets. In fact, `a[n]` is really just syntactic sugar for `*(a + n)`. Remember that pointers are a different type: when you add an integer to a pointer, it's scaled by something akin to the `sizeof` function.

Some examples:

- `int a[5];` declares an array of five `ints` (on Zeus this is 20 bytes)
- `a[0]` references the value at address `a`
- `a[1]` references the value at the address `(a+1)`.

- As noted above, in hardware this would be the address of `a` plus four bytes (since an `int` on Zeus is four bytes)
- `a[4]` references the last element of `a`
- `a[5]` is out of bounds but completely legal C – no bounds checking, remember?

Arrays and Pointers

Files: `ary_a.c` and `ary_b.c`.

The name of the array is the address at of the start of the array. An array name is *not* a variable – it is a constant value.

This may be surprising. “But an array name is an `lvalue`!” you might claim – and you’d be correct. However, an array name is not a *modifiable lvalue*. For an excellent breakdown, see this Stack Overflow post:

Is there a reason why an array name is not an lvalue?

Pointer Arithmetic in Arrays

File: `ptr_ary.c`.

You can use pointer arithmetic when working with arrays.

```
1  int nums[5] = {1, 6, 10, 42, -14};
2  int *p_nums = NULL;
3  p_nums = nums;
4  // Address of the array
5  printf("Address of nums: %p\n", nums);
6  printf("Address p_nums points to: %p\n", p_nums);
7  // Access the fourth number with [] notation
8  printf("p_nums[3] == %d\n", p_nums[3]);
9  // Access the fourth number with pointer arithmetic
10 printf("p_nums[3] == %d\n", *(p_nums + 3));
```

Question 3

Given the following snippet:

```
1  int nums[5] = {1, 6, 10, 42, -14};
2  int *p_nums = &nums[3];
3  p_nums = nums;
4  // Address of nums = 0x400
5  // sizeof(int) = 4 bytes, sizeof(int *) = 8 bytes
```

what is the value of each of the following expressions?

1. `&nums[2]`
 - 0x408
2. `p_nums[-1]`
 - 10
3. `*(p_nums + 1)`
 - -14

I/O in C

What is `stdio.h`?

- The header file for the C Standard Library (`libc.a`)

Okay, so what's `libc.a`?

- A static library which is automatically included
- It's a set of pre-compiled C objects which contain all of the core C functions
- `stdio.h` only has the prototypes and macros for those functions

Some examples of standard I/O operations:

- Opening and closing files: `fopen` and `fclose`
- Reading and writing files: `fread` and `fwrite`
- Reading and writing text: `fgets` and `fputs`
- Formatted reading and writing: `fscanf` and `fprintf`

Standard I/O is implemented on top of the Operating System (OS) I/O:

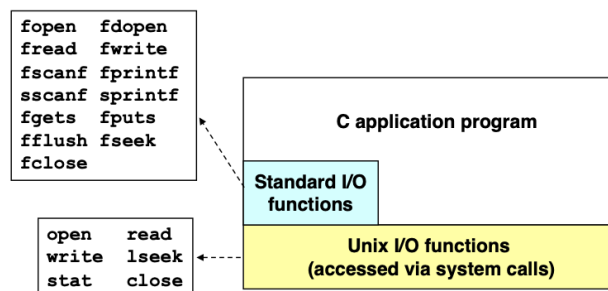


Figure 2: Unix I/O interacting with C I/O.

Standard I/O Streams in C and Unix

Standard I/O models all open files as streams. This is an abstraction for a file descriptor and a buffer.

There are three streams all C programs start with access to:

Name	Type	Buffered?	File Descriptor
<code>stdin</code>	Standard Input	yes	0
<code>stdout</code>	Standard Output	yes	1
<code>stderr</code>	Standard Error	no	2

Common I/O Functions

These are all declared within the `stdio.h` header.

Function	Description
<code>putchar</code>	Displays an ASCII character to the screen
<code>getchar</code>	Reads an ASCII character from the keyboard
<code>printf</code>	Displays a formatted string
<code>scanf</code>	Reads a formatted string
<code>fopen</code>	Open or create a file for I/O
<code>fprintf</code>	Writes a formatted string to a file
<code>fscanf</code>	Reads a formatted string from a file

Formatted I/O

Both `printf` and `scanf` allow conversion between ASCII representations and internal data types.

The format string of these functions contains text to be read/written as well as the format characters which describe their formatting.

Code	Description
<code>%d</code>	Signed decimal integer
<code>%f</code>	Signed decimal floating-point number

Code	Description
<code>%x</code>	Hexadecimal number
<code>%e</code>	Scientific notation
<code>%c</code>	ASCII character
<code>%s</code>	ASCII string
<code>%p</code>	Memory address (hex)

Both `printf` and `scanf` allow special characters as well.

Code	Description
<code>\n</code>	Newline
<code>\t</code>	Tab
<code>\b</code>	Backspace
<code>\\</code>	Backslash
<code>\'</code>	Single quote
<code>\"</code>	Double quote
<code>\0nnn</code>	ASCII code <i>nnn</i> in octal
<code>\xnnn</code>	ASCII code <i>nnn</i> in hex

printf

File: `print.c`.

Prints its first argument (the format string) to `stdout` with all of the formatting characters replaced by the ASCII representation of the corresponding arguments.

As an example, the snippet

```
1 int a = 100;
2 char c = 'z';
3 char hw[13] = "Hello World!";
4 double pi = 3.14159;
5 printf("a: %d, \'%c\', 0x%x\n", a, a, a);
6 printf("c: %d, \'%c\', 0x%x\n", c, c, c); printf("hw: \"%s\"\n", hw);
7 printf("pi: %lf, %.2lf, %e\n", pi, pi, pi);
```

prints

```
1 $ ./print
2 a: 100, 'd', 0x64
3 c: 122, 'z', 0x7a
4 hw: "Hello World!"
5 pi: 3.141590, 3.14, 3.141590e+00
```

when executed.

printf Pitfalls

What happens when you forget a data argument, like in the statement `printf("a: %d, \\'%c\\', 0x%x\\n", a, a);`?

Each code (`%d`, `%x`, etc.) will convert whatever is on the stack (or register) where it expects that argument to be located. So you'll get *some* value... perhaps just not the one you wanted.

Takeaway: *every register and each byte of RAM has a value.*

scanf Conversions

For each conversion, `scanf` will skip whitespace characters and then read ASCII characters until it encounters the first character that should *not* be included in the converted value.

Code	Stops when
<code>%d</code>	Reads first non-digit
<code>%x</code>	Reads first non-digit
<code>%s</code>	Reads until first whitespace character

Data arguments *must be pointers* because `scanf` stores the converted value at the given memory address.

scanf

File: `scan.c`.

Reads ASCII characters from `stdin`, then:

1. Matches characters to the next expected code in the format string
2. Converts the matched characters according to the format

3. Stores the converted values in the provided address
 - Repeats until all codes have been read
4. Returns the number of successful conversions

scanf Question

Consider the following snippet:

```
1 char name[100];
2 int age, gnum;
3 double gpa;
4 scanf("%s %d/%d %lf", name, &age, &gnum, &gpa);
5 printf("%s %d/%d %lf\n", name, age, gnum, gpa);
```

1. What will be printed out with this input: Kevin 73/123456 3.92

```
1 $ ./scan
2 Kevin 73/123456 3.92
3 Kevin 73/123456 3.92
```

2. What will be printed out with this input: Kevin 73 123456 3.92

```
1 $ ./scan
2 Kevin 73/123456 3.92
3 Kevin 73/0 0.00
```

fgets and sscanf

These are safer ways to get user input.

It's better to read the entire input as a string using `fgets` and then parse it with `sscanf`.

Dynamic Allocation

The standard C library provides a function for allocating memory at run-time to the heap. The function prototype is `void *malloc(int numBytes);`.

It returns a generic pointer `void *` to a contiguous region of memory of the requested size (in bytes).

The bytes are allocated from a region in memory called the heap.

- The OS keeps track of chunks of memory from the heap that have been allocated
- We'll be studying this system in depth this semester

The standard C library also provides a function for deallocating memory at run-time from the heap. The function prototype is `void *free(void *)`;

If data allocated by `malloc` is not freed, then over the course of your program, you could run out of heap memory.

- This is called a *memory leak*
- If this occurs, your program may crash

Data Structures

A data structure is just a particular organization of data in memory

- Group related items together
- Organize data to be efficient to execute and convenient to program

Data structures we'll examine in this class:

- Arrays: contiguous, in memory, homogenous data types (special case product type)
- Structs: grouped heterogeneous data types (product type)
- Linked Lists: non-contiguous, dynamically allocated type (sum type)

Structs

A `struct` lets you group different data together.

As an example, let's represent a wireless packet containing drone data from flight:

```
1 struct flight_type_t {
2     char flight_num[7];
3     int altitude; // In meters
4     int longitude; // In hundredths
5     int latitude; // In hundredths
6     int heading; // In radians
7     double speed; // In m/s
8 };
```

A `struct` definition does not allocate memory.

To allocate memory, we need to declare a variable:

```
1 struct flight_type_t drone_one;
2 struct flight_type_t drone_two;
3 struct flight_type_t drone_three;
4 drone_one.altitude = 10000;
5 drone_one.longitude = -7730;
```

Our stack might then look like (where addresses near the top are higher, since this is a stack):

Stack	Variable
:	:
-7730	drone_one.longitude
10000	drone_one.altitude
	drone_one.flight_num[7]
:	:
	drone_one.flight_num[0]

typedef

A `typedef` lets you declare a type synonym.

The syntax is `typedef <type> <name>;`. As an example, we can make a `typedef` for the drone data.

```
1 typedef struct flight_type_t {
2     char flight_num[7];
3     int altitude; // In meters
4     int longitude; // In hundredths
5     int latitude; // In hundredths
6     int heading; // In radians
7     double speed; // In m/s
8 } FlightType;
```

Declaring variables would then look like

```
1 struct FlightType drone_one;
2 struct FlightType drone_two;
3 struct FlightType drone_three;
4 drone_one.altitude = 10000;
5 drone_one.longitude = -7730;
```

Notice that the use of the dot operator is unchanged.

Dynamic Allocation of Structs

Generally, you'll put your struct allocations on the heap.

```
1 // Some kind of setup
2 FlightType *drone_four = NULL;
3 drone_four = (FlightType *) malloc(sizeof(FLIGHT_TYPE)); // Cast the pointer!
4 strncpy(drone_four->flight_num, "DRN4", 4); // Remember > for dynamic alloc
   structs drone_four->altitude = 10231;
```

```
5 // Some kind of use of the data  
6 free(drone_four);
```