

---

# Machine Level Representation IV

x86-64 Assembly Data Structures

Connor Baker



Department of  
**Computer Science**

Compiled on October 27, 2019 at 6:28pm

## Reading Track

This Lecture covers Chapters 3.8-3.9 (Assembly Expressions)

We'll spend several classes on each of the Assembly Lectures.

Reading Homework:

1. Read through Chapter 3.10-3.11 (Stack and Buffer Overflows)
2. Continue to Review Chapters 3.1-3.9 This semester we'll be working with all of these topics in Assembly
  - Data Types and Sizes in x86-64 Assembly (Arrays of different types!)
  - Moving Data and Arithmetic
  - Control Flow (If/Else, Loops, and Switches)
  - Procedures (Callee-Save/Stack Save) with and without Recursion
  - One and Two dimensional Arrays, Structs, and Linked Lists


## Lecture Overview

ASM Structures	Implement Data Structures in x86-64 Assembly	Describe Arrays in Memory	3.8
		Access an Element from a Single-Dimensional Array in Assembly	
		Access an Element from a Multi-Dimensional Array in Assembly	3.9
		Describe Structs in Memory	
		Describe Unions in Memory	
		Describe Alignment of Data Structures in Memory	

**Figure 1:** Lecture overview.

## Array Allocation


The basic principle: `T A[N]` declares an array of data type `T` for integer constant `N` elements. The size of the array is `N * sizeof(T)` bytes.



**Figure 2:** Sample array allocations.

**Question 1**

The identifier A is the address of the first element of the array. This address can also be used as a pointer to the array element 0.



**Figure 3:** An example array.

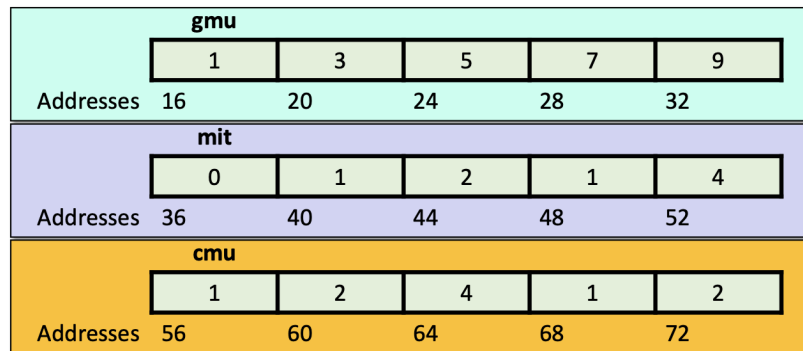
C Code	Type	Value
A[4]	int	8
A	int *	x
A+1	int *	x+4
&A[3]	int *	x+12
*(A+3)	int	-3

**Array Allocation Example**

Assume that we have the following arrays initialized:

```
int gmu[5] = {1, 3, 5, 7, 9};
int mit[5] = {0, 1, 2, 1, 4};
int cmu[5] = {1, 2, 4, 1, 2};
```

Each array's elements are contiguous in Memory. Each of the arrays may not be in successive 20-byte blocks!



**Figure 4:** Array contents in memory.

## Accessing an Array Element

File: `array_access.s`

The starting address for an element in a array is  $A + i * K$  where  $A$  is the base address of the array,  $i$  is the index, and  $K$  is the data type size.

In Assembly, we need to use one of the forms ot access this element.

```
# %rdi = gmu
# %rsi = i
# Put the result in eax
movl (%rdi,%rsi,4), %eax # gmu[i]
```

File: `array_loop longs.c`

Let's write some assembly to sum up all the values of the array.

```
long sum(int *ary, long len) {
    long ans = 0;
    for(long i = 0; i < len; i++) {
        ans += ary[i];
    }
    return ans;
}
int main() {
    int gmu[5] = {1,3,5,7,9};
    return sum(gmu, 5);
}
```

The corresponding assembly might look like

```
sum:                                #rdi has ary and rsi has number of elements in ary
    movq $0, %rcx                   # long i = 0
    movq $0, %rax                    # long ans = 0
```

```

    jmp test                # Jump to Middle style For loop
top:
    addl (%rdi, %rcx, 4), %eax # addl because ary is an int array
    incq %rcx                # ++i;
test:
    cmpq %rsi, %rcx         # compare i against len
    jl top
    ret                    # returns 64-bit ans in %rax

```

But what if we changed the return type of `sum` to `int`?

```

long sum(int *ary, int len) {
    int ans = 0;
    for(int i = 0; i < len; i++) {
        ans += ary[i];
    }
    return ans;
}
int main() {
    int gmu[5] = {1,3,5,7,9};
    return sum(gmu, 5);
}

```

We would need to modify our assembly:

```

sum:
    movl $0, %ecx          # rdi has ary and esi has number of elements in ary
    movl $0, %eax          # int i = 0
    jmp test              # int ans = 0
top:
    movslq %ecx, %rcx      # Jump to Middle style For loop
    addl (%rdi, %rcx, 4), %eax # Sign extend i to 64-bits for pointer math
    incl %ecx              # addl because ary is an int array
test:
    cmpl %esi,%ecx         # ++i;
    jl top                # compare i against len
    ret                    # returns 32-bit ans in %eax

```

## Digression – Data Type Errors in Assembly

Consider

```
addl (%rdi, %ecx, 4), %eax # Bad Code
```

That would yield the following error message:

```
array_loop_ints.s:23: Error: '(%rdi,%ecx,4)' is not a valid base/index expression
```

For dereferencing forms, this is address arithmetic, so it should be 64-bit register, like as follows:

```
addl (%rdi, %rcx, 4), %eax # Good Code
```

Again, if we used an incorrect suffix

```
incq %ecx # Bad Code
```

we would get the error message

```
array_loop_ints.s:24: Error: incorrect register '%ecx' used with 'q' suffix
```

This is a 64-bit operation (incq) on a 32-bit register name (ecx). Types have to match!

```
incl %ecx # Good Code
```

## Nested Arrays

Nested arrays take the form of `T D[R][C]`. This form of a nested array has an array of `R` elements (rows), each of which has an entry of `C` elements (columns). In memory, the size of this array is  $R * C * K$  bytes, where `K` is the size of data type `T`.

As an example, `int A[10][5]` would take  $10 * 5 * 4 = 200$  bytes.

## Nested Array Example

```
int ary[4][5] = {{1, 5, 2, 0, 6},
                 {1, 5, 2, 1, 3 },
                 {1, 5, 2, 1, 7 },
                 {1, 5, 2, 2, 1 }};
```

The array `ary` is an array of four elements, allocated contiguously. Each element of `ary` is an array of 5 `ints`, allocated contiguously.

Nested arrays are ordered in memory in Row-Major ordering. This is analogous to all elements of the zeroth row being followed by the elements of the first row, and so on.

## Row Vectors and Accessing Elements

Consider the following array:

1	5	2	0	6	1	5	2	1	3	1	5	2	1	7	1	5	2	2	1
76	80	84	88	92	96	100	104	108	112	116	120	124	128	132	136	140	144	148	152

**Figure 5:** The array `int ary[4][5]`.

To access the element `ary[i][j]`, we use the formula  $A + (i * C * K) + (j * K)$  where  $A$  is the base address,  $i$  is the  $i$ th row vector,  $C$  is the number of columns,  $j$  is the  $j$ th column vector, and  $K$  is the size of the data type.

We can use the simplified mathematical version

$$A + (i \times C \times K) + (j \times K) \equiv A + (i \times C + j) \times K.$$

To access some arbitrary element of the array `ary` where `%rdi` holds  $i$  and `%rsi` holds  $j$ , we would use the following assembly (remember that `ary` is a static defined region, so you can use it like an immediate address)

```
leaq (, %rsi, 4), %rdx      # K * j
leaq (%rdi, %rdi, 4), %rdi  # C * i
movl ary(%rdx, %rdi, 4), %eax #*(ary+i*C*+j*K)
```

(The second instruction sets `%rdi` to five times the previous value since `ary` is a  $4 \times 5$  array.)

## Array Practice

### Question 3

Suppose we have the following arrays that we will store sequentially starting at memory address `0x100`.

```
int  A[20];
char B[40];
int  C[200];
```

For your convenience, the following conversions are provided.

Math Hints

- $40 = 0x28$
- $80 = 0x50$
- $800 = 0x320$

Array	Bytes Needed	Base Address	Equation for <code>arr[i]</code>
<code>A[]</code>	$20 * 4 = 80$ ( $0x50$ )	<code>0x100</code>	$0x100 + i * 4$
<code>B[]</code>	$40 * 1 = 40$ ( $0x28$ )	<code>0x150</code>	$0x150 + i * 1$
<code>C[]</code>	$200 * 4 = 800$ ( $0x320$ )	<code>0x178</code>	$0x178 + i * 4$

## Nested Array Practice

### Question 4

Suppose we have the following arrays that we will store sequentially starting at memory address `0x1000`.

```
int A[10][10];
char B[40][5];
int C[20][10];
```

For your convenience, the following conversions are provided.

Math Hints

- `200 = 0xC8`
- `400 = 0x190`
- `800 = 0x320`
- `1600 = 0x640`

Array	Bytes Needed	Base Address	Equation for <code>arr[i][j]</code>
<code>A[][]</code>	<code>10 * 10 * 4 = 400 (0x190)</code>	<code>0x1000</code>	<code>0x1000 + (i * 10 * 4) + (j * 4)</code>
<code>B[][]</code>	<code>40 * 5 * 1 = 200 (0xC8)</code>	<code>0x1190</code>	<code>0x1190 + (i * 5 * 1) + (j * 1)</code>
<code>C[][]</code>	<code>20 * 10 * 4 = 800 (0x320)</code>	<code>0x1258</code>	<code>0x1258 + (i * 10 * 4) + (j * 4)</code>

## Multi-Level Arrays


A multi-level array is an array of pointers to other arrays.

For example, if we had the following arrays defined

```
int gmu[5] = {9, 4, 7, 2, 0};
int mit[5] = {0, 2, 1, 3, 9};
int cmu[5] = {1, 5, 2, 1, 3};
int *univ[3] = {mit, cmu, gmu};
```

then in memory they might look like





**Figure 6:** An example of a multi-level array.

## Accessing a Multi-Dimensional Array Element

Our goal is to retrieve `univ[i][j]` (as defined previously).

We start by acquiring the address of the array that `univ` is referencing. The address is retrieved by dereferencing `univ + 8 * i`.

Then, we must index into that retrieved array: `(univ + 8 * i) + 4 * j`.

With assembly (assuming `%rdi` holds `i` and `%rsi` holds `j`)

```
leaq (, %rsi, 4), %rsi    # 4*j
addq univ(, %rdi, 8), %rsi # M[univ+8*i]+4*j
movl (%rsi), %eax        # M[M[univ+8*i]+4*j]
```

## Fixed-Size Arrays

If the dimensions of the arrays are constant, you can optimize the code further by pre-computing the offsets.

As an example, suppose that we have a two-dimensional,  $16 \times 16$  array, `int A[16][16]`. Then

$$\begin{aligned}
 \text{int } x &= A[i][j] \\
 &= M[A + (i \times C \times K) + (j \times K)] \\
 &= M[A + 64i + 4j]
 \end{aligned}$$

Our assembly would change thusly:

```
# a in %rdi, i in %rsi, j in %rdx
salq $6, %rsi    # 64i
addq %rsi, %rdi  # a + 64i
movl (%rdi,%rdx,4), %eax # M[a + 64i + 4j]
ret
```


## Structs

C `structs` are product types which allow for a heterogeneous collection of types. They are represented by a block of memory in which all fields appear in order of declaration. The compiler determines the overall size.

For example, suppose that we have the following `struct`:


```
struct rec {
    int a[4];
    long i;
    struct rec *next;
};
```

With C, our in-memory representation of the struct might look like



**Figure 7:** An example of the `rec struct` in memory as viewed by C.

With assembly, given that `%rdi` contains the address of an instance of the `struct`, our memory would look like



**Figure 8:** An example of the `rec struct` in memory, where `%rdi` is the base of the `struct`.


The offset of each structure member must be determined at compile time.

Assume that we want to get the address of `r->a[idx]` using assembly.

```
# r in %rdi, idx in %rsi
leal (%rdi,%rsi,4), %eax
ret
```

## Linked Lists

Assuming the same `struct rec`, how would we get `r->next->next->i`?



**Figure 9:** A linked list.

```
# r starts out with the value 0
# Hint: Start with r = r->next;
movq 24(%rdi), %rdi # r = r->next = 100
movq 24(%rdi), %rdi # r = r->next = 40
movq 16(%rdi), %rax # x = r->i = 150
```

## Struct Alignment

Systems may use alignment restrictions for efficiency. Alignment is making sure each data type is organized such that its address starts on a multiple of its own size.

For Example: an `int` is 4 bytes, so it starts on an address that is a multiple of 4 bytes, like `0x600400` or `0x0400698`.

For alignment, an Array should be aligned to the data type inside.



For alignment, a `struct` should be aligned to the biggest data type inside. On this system, we will only be looking at alignment for `structs`.

For the following two rules assume the following struct:

```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```

**Rule One**



All data types in a `struct` must start on a multiple of that type size.



**Figure 10:** The first rule of `struct` alignment.

**Rule Two**

The `struct` size must be a multiple of the largest type inside of it.



**Figure 11:** The second rule of `struct` alignment.

**Struct Optimal Ordering**


To reduce wasted space, reorder the structs to put the largest types first.

Compare the result of the first struct

```

struct S1 {
    char c;
    double v;
    char x;
} *p;

```




**Figure 12:** A mal-ordered `struct`.

with that of the organized struct

```

struct S1 {
    double v;
    char c;
    char x;
} *p;

```



**Figure 13:** A well-ordered `struct`.


## Array of Structs

Since each `struct` is aligned to some multiple of the largest data type arrays of `structs` are easy to handle.

```


struct S2 {
    double v;
    int i[2];
    char c;
} *p;

```



**Figure 14:** The layout of the above `struct` in memory.

An array of such `struct` might look like



**Figure 15:** An array of `struct` in memory..

## Unions


C unions are data types to allocate memory for the largest defined element.

They essentially allow for reinterpretation of data.

A `struct` like

```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```

looks like




**Figure 16:** The `struct S1` in memory.

The union

```
union U1 {
```

```
char c;  
int i[2];  
double v;  
} *p;
```

looks like



**Figure 17:** The `union U1` in memory.

Note that the `union` only allocates memory equivalent to the size of the largest member.